

Lecture 2

CS5487 Lecture Notes (2022B)
Prof. Antoni B. Chan
Dept of Computer Science
City University of Hong Kong

Parameter Estimation

How do we find a prob. dist for a r.v. X ?

Three Steps:

1) Choose a parametric model (e.g. Gaussian)
 Θ = parameters.

2) collect samples from r.v. X :

$$D = \{x_1, \dots, x_N\}$$

We assume x_i 's are independent; x_i are iid samples
independent & identically distributed

3) Maximum likelihood principle:

the optimal parameter Θ^* is that which maximizes the probability (likelihood) of the training data.

$$\Theta^* = \underset{\Theta}{\operatorname{argmax}} p(D|\Theta)$$

*likelihood of data w.r.t. param Θ .
"likelihood function"*

$$= \underset{\Theta}{\operatorname{argmax}} \log p(D|\Theta)$$

\log = log-likelihood function.

$$= \underset{\Theta}{\operatorname{argmin}} -\log p(D|\Theta)$$

negative LL function (loss)

Note: D is known, so $p(D|\Theta)$ is a function of Θ .
It is not a probability w.r.t. Θ .

\log = natural \log (log base e)

Data LL

$$l(\Theta) = \log p(D|\Theta)$$

independence assumption

$$= \log \prod_{i=1}^N p(x_i|\Theta)$$

$$l(\Theta) = \sum_{i=1}^N \log p(x_i|\Theta)$$

To get the MLE solution


if Θ is a scalar, at local optimum:

$$1) \frac{\partial}{\partial \Theta} \log p(D|\Theta) = 0 \quad \text{at } \Theta^*$$

$$2) \frac{\partial^2}{\partial \Theta^2} \log p(D|\Theta) < 0 \quad \text{at } \Theta^*$$

(local maximum; concave)

3) check the boundary conditions of Θ (if necessary)



if Θ is a vector:

$$1) \nabla_{\Theta} l(\Theta) = \begin{bmatrix} \frac{\partial}{\partial \Theta_1} l(\Theta) \\ \vdots \\ \frac{\partial}{\partial \Theta_p} l(\Theta) \end{bmatrix} = 0$$

gradient



$$2) \nabla_{\Theta}^2 l(\Theta) = \begin{bmatrix} \frac{\partial^2}{\partial \Theta_1^2} & \dots & \frac{\partial^2}{\partial \Theta_1 \partial \Theta_p} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2}{\partial \Theta_p \partial \Theta_1} & \dots & \frac{\partial^2}{\partial \Theta_p^2} \end{bmatrix} l(\Theta)$$

Hessian

(negative definite)

$H < 0$: negative definite: $\Theta^T H \Theta < 0, \forall \Theta$

$H > 0$: positive defn: $\Theta^T H \Theta > 0, \forall \Theta$

"mountain" concave in all dir.
"bowl" - convex in all directions.

Example: Bernoulli

$$\theta = \pi, \quad 0 \leq \pi \leq 1, \quad x = \{0, 1\}$$

$$\begin{aligned} \ell(\theta) &= \sum_{i=1}^N \log p(x_i | \theta) \\ &= \sum_{i=1}^N \log [\pi^{x_i} (1-\pi)^{(1-x_i)}] \\ &= \sum_i [x_i \log \pi + (1-x_i) \log (1-\pi)] \\ &= \underbrace{\left(\sum_i x_i \right) \log \pi}_{\# \text{ of } 1\text{'s}} + \underbrace{\left(\sum_i (1-x_i) \right) \log (1-\pi)}_{\# \text{ of } 0\text{'s}} \end{aligned}$$

$m = \sum_i x_i \leftarrow$ "sufficient statistic" = $\ell(\theta)$ only depends on the N observations (dataset) through this value.

$$\ell(\theta) = m \log \pi + (N-m) \log (1-\pi)$$

find the max:

$$1) \frac{\partial}{\partial \pi} \ell(\theta) = \frac{m}{\pi} + \frac{N-m}{1-\pi} (-1) = 0 \quad \downarrow \times \pi(1-\pi)$$

$$(1-\pi)m - \pi(N-m) = 0$$

$$m - m\pi - N\pi + m\pi = 0$$

$$m - N\pi = 0 \Rightarrow \boxed{\hat{\pi} = \frac{m}{N} = \frac{1}{N} \sum_{i=1}^N x_i}$$

"fraction of 1's observed"
(sample mean)

$$\begin{aligned} 2) \frac{\partial^2}{\partial \pi^2} \ell(\theta) &= \frac{\partial}{\partial \pi} \left(\frac{\partial}{\partial \pi} \ell(\theta) \right) = \frac{\partial}{\partial \pi} \left(\frac{m}{\pi} - \frac{N-m}{1-\pi} \right) \\ &= \frac{-m}{\pi^2} - \frac{N-m}{(1-\pi)^2} (-1)(-1) < 0 \quad \checkmark \end{aligned}$$

$$3) \text{ boundary condition: } 0 \leq m \leq N \quad 0 \leq \frac{m}{N} \leq 1 \quad \checkmark$$

Example: Gaussian

$$① \theta = \mu \quad (\sigma^2 \text{ known})$$

$$\begin{aligned} \ell(\theta) &= \sum_{i=1}^N \log p(x_i | \theta) \\ &= \sum_i \left[-\frac{1}{2} \log 2\pi - \frac{1}{2} \log \sigma^2 - \frac{1}{2\sigma^2} (x_i - \mu)^2 \right] \\ &= -\frac{N}{2} \log 2\pi - \frac{N}{2} \log \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2 \end{aligned}$$

what are the sufficient statistics?
 $\left\{ \sum_i x_i, \sum_i x_i^2 \right\}$

max wrt μ :

$$\frac{\partial \ell(\theta)}{\partial \mu} = -\frac{1}{2\sigma^2} \sum_i 2(x_i - \mu)(-1) = 0$$

$$\sum_i (x_i - \mu) = 0 \Rightarrow \sum_i x_i - N\mu = 0 \Rightarrow \boxed{\hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i}$$

(sample mean)

$$② \theta = \sigma^2 \quad (\mu \text{ is known})$$

$$\begin{aligned} \frac{\partial \ell(\theta)}{\partial \sigma^2} &= -\frac{N}{2} \frac{1}{\sigma^2} - \frac{1}{2\sigma^4} (-1) \sum_i (x_i - \mu)^2 = 0 \quad \downarrow \times \sigma^4 \\ &= -\frac{N}{2} \sigma^2 + \frac{1}{2} \sum_i (x_i - \mu)^2 = 0 \end{aligned}$$

$$\boxed{\hat{\sigma}^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2} \quad \text{"Sample variance"}$$

M.V. Gaussian

$$\begin{cases} \hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i \\ \hat{\Sigma} = \frac{1}{N} \sum_{i=1}^N (x_i - \hat{\mu})(x_i - \hat{\mu})^T \end{cases}$$

See next tutorial...

Estimators

the Estimate (e.g. $\hat{\mu}$) is a number.

the Estimator is a r.v. (over possible datasets)

$$\text{estimator } f(X_1, \dots, X_N) = \frac{1}{N} \sum_{i=1}^N X_i$$

\uparrow r.v. for each sample
 $X_i \sim p(X_i|\theta)$ + true distribution.

The estimate is the value of the estimator for a given dataset D .

$$\hat{\mu} = f(X_1, \dots, X_N) \Big|_{X_i = x_i, \dots} = \frac{1}{N} \sum_{i=1}^N x_i$$

\uparrow sample

Since the estimator is a r.v., we can derive the mean & variance to quantify the "goodness".

Bias & Variance

 $\hat{\theta} = f(X_1, \dots, X_N)$

1) Will it converge to the true value of θ ?

$$\text{Bias}(\hat{\theta}) = E_{X_1, \dots, X_N} [\hat{\theta} - \theta] = E_X [\hat{\theta}] - \theta$$

\uparrow true value \uparrow mean of the estimator.

if the bias is non-zero, then we can never get the true value (even if infinite samples).

2) How long will it take to converge? (How many samples do we need?)

$$\text{var}(\hat{\theta}) = E_{X_1, \dots, X_N} [(\hat{\theta} - E\hat{\theta})^2]$$

Example: Gaussian

$$\text{Estimator: } \hat{\mu} = \frac{1}{N} \sum_{i=1}^N X_i$$

$$\text{Mean of } \hat{\mu}: E_{X_1, \dots, X_N} \left[\frac{1}{N} \sum_{i=1}^N X_i \right] = \frac{1}{N} \sum_{i=1}^N E_{X_i} [X_i] = \frac{1}{N} \cdot N \mu = \mu$$

\uparrow mean of true distribution (μ)

Bias of $\hat{\mu} = 0$ ✓

$$\text{Var of } \hat{\mu}: E_{X_1, \dots, X_N} [(\hat{\mu} - E\hat{\mu})^2] = E \left[\left(\frac{1}{N} \sum_{i=1}^N X_i - \mu \right)^2 \right]$$

$\left(\frac{1}{N} \sum_{i=1}^N (X_i - \mu) \right)^2$

$$= \frac{1}{N^2} E \left[\left(\sum_{i=1}^N (X_i - \mu) \right)^2 \right]$$

$(a+b)^2 = a^2 + 2ab + b^2$

$$= \frac{1}{N^2} E \left[\sum_{i=1}^N \sum_{j=1}^N (X_i - \mu)(X_j - \mu) \right]$$


$\left(\sum_{i=1}^N a_i \right)^2 = \sum_{i,j} a_i a_j$

$$i=j \Rightarrow E[(X_i - \mu)^2] = \sigma^2$$

$$i \neq j \Rightarrow E[(X_i - \mu)(X_j - \mu)] = 0$$

$$= \frac{1}{N^2} (N \sigma^2) = \frac{\sigma^2}{N} = \text{var}(\hat{\mu})$$

\uparrow variance converges to 0 as $N \rightarrow \infty$.



Gaussian variance (PS 2-12)

$$E(\hat{\sigma}^2) = \frac{N-1}{N} \sigma^2 \Rightarrow \text{Bias}(\hat{\sigma}^2) = \underbrace{-\frac{1}{N} \sigma^2}_{\text{true variance}} \neq 0$$

to make it unbiased:

$$\hat{\hat{\sigma}}^2 = \frac{N}{N-1} \hat{\sigma}^2 = \frac{N}{N-1} \frac{1}{N} \sum_i (x_i - \mu)^2 = \frac{1}{N-1} \sum_i (x_i - \mu)^2$$


Important Asymptotic Properties of MLE

1) consistent - As $N \rightarrow \infty$, the estimate converges to the true value. Asymptotically unbiased.

2) efficient - achieves the Cramér-Rao Lower Bound (CRLB) as $N \rightarrow \infty$.

- CRLB is a theoretical bound on the variance of any unbiased estimator for a given $p(x|\theta)$.
- i.e. no unbiased estimator can get lower variance.

MLE for Regression



$x \in \mathbb{R}$ input
 $y \in \mathbb{R}$ output
 learn $f(x)$

consider a polynomial function (k^{th} order)

$$f(x, \theta) = \sum_{d=0}^k x^d \theta_d = \underbrace{\begin{bmatrix} 1 \\ x \\ x^2 \\ \vdots \\ x^k \end{bmatrix}}_{\phi(x)}^T \underbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_k \end{bmatrix}}_{\theta} = \underbrace{\phi(x)^T \theta}_{\text{linear function in } \theta}$$

observe a noisy output:

$$y = f(x, \theta) + \epsilon \quad \leftarrow \text{noise } \epsilon \sim \mathcal{N}(0, \sigma^2), \text{ i.i.d.}$$

equivalently, (y is a r.v.)

$$p(y|x, \theta) = \mathcal{N}(y | f(x, \theta), \sigma^2)$$

Given dataset $\{(x_i, y_i)\}_{i=1}^N$, estimate θ using MLE:

$$\begin{aligned} \hat{\theta} &= \underset{\theta}{\operatorname{argmax}} \sum_i \log p(y_i | x_i, \theta) \\ &= \underset{\theta}{\operatorname{argmin}} \sum_i (y_i - f(x_i, \theta))^2 \quad \leftarrow \text{least-squares formulation} \\ &= \underset{\theta}{\operatorname{argmin}} \|y - \Phi^T \theta\|^2, \quad \Phi = [\phi(x_1) \dots \phi(x_N)], \quad y = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix} \end{aligned}$$

$$\hat{\theta} = (\Phi \Phi^T)^{-1} \Phi y$$

Notes:

- 1) MLE is more general than LS.
- 2) Assumptions are explicit
 - i) Gaussian noise
 - ii) $\mu=0$, σ^2 variance (fixed)
 - iii) noise is iid.
- 3) MLE can describe other LS formulations:
 - i) weighted LS (PS 2.8)
 - ii) regularized LS (lecture 3)
 - iii) L_p -norm (PS 2.9)

generalized linear models (GLM)